

RESEARCH ROADMAP

STATUS UPDATE - 2026-10-02

How to read this update. This addendum supplements the original 22-page roadmap. It records work completed in the repository, generated evidence, and remaining dependencies. Earlier archived results remain historical records and are not silently replaced. All current outputs are simulator-derived; none is field validation or a QKD security proof.

1. Work completed in this update

Workstream	Change / result	Remarks and evidence
Corrected route comparison	Baseline selectors use the observation's feasible-action ordering; all methods use the same season and channel seeds. Per-episode output records action validity, no-feasible-action states, fallbacks, route, hops, revisits, and link validity.	The comparison output is identified below. No-feasible-action states are separated from invalid policy actions.
Physics consistency	Independent formula transcription covered 201 fiber distances from 0 to 200 km. Maximum absolute difference: QBER 0; normalized proxy 0. Transmission and proxy decrease; QBER does not decrease.	Internal consistency only. It does not validate assumptions against measurements or establish security.
FSO stochastic checks	All 16 configured season/distance anchors sampled at 2,000 trials each; maximum absolute availability deviation 0.0125 (1.22 target standard errors). Log-Cn2 AR(1): configured 0.85, sampled 0.8522 over 20,000 steps.	These are model targets and samples, not validation against Indian weather observations.
Link-feature sensitivity	The controlled diagnostic includes a proxy-rate ratio sweep, a feasible QBER-margin sweep, and an outage-mask counterfactual with route context held fixed. The FSO action probability rose from 0.311 to 0.638 as its proxy-rate ratio rose from 0.25x to 4x; at QBER 0.111 it was masked. In a separate 32-sample FSO turbulence counterfactual, a 100x Cn2 multiplier lowered median proxy from 0.202 to 0.183 and mean FSO action probability from 0.749 to 0.740; FSO remained selected in 100% of samples.	A policy-input diagnostic does not establish link-measurement validity or route-level performance.
Five-seed stability campaign	1/5 independent policies have final evaluations. Seed 20261003 is the latest run without a final summary. Completed seed 20261002: 93.8% success, 69.80 mean hops among successful episodes.	Inter-seed spread and confidence summaries are withheld until all planned seeds finish. A training-seed interval is distinct from an episode-level interval.

2. Paired route comparison

Policy	Success	Successful mean hops	Invalid actions	No-valid-action states	Revisits
GNN-PPO	16/16 (100.0%; 95% CI 80.6%-100.0%)	28.25	0	0	2
Random	0/16 (0.0%; 95% CI 0.0%-19.4%)	n/a	0	8	2459
Dijkstra-km	4/16 (25.0%; 95% CI 10.2%-49.5%)	233.00	0	0	172
BFS-hop	16/16 (100.0%; 95% CI 80.6%-100.0%)	28.25	0	0	3
Max-SKR	0/16 (0.0%; 95% CI 0.0%-19.4%)	n/a	0	5	984

Conditions: Delhi to Chennai; 4 shared environment seeds per each of 4 seasons; QBER/rate action mask; 400 step limit. Scope is single-route simulator behavior, not traffic blocking, network-wide key-pool balance, or delivered service capacity.

Checkpoint transfer across city pairs

The same 50-epoch checkpoint completed 31/42 ordered city-pair episodes (73.8%) under the archived monsoon-night scenario. Median successful route length was 20 hops; localized decision latency was 3.83 ms median and 4.34 ms p95 on this host. This is one

stochastic sample per ordered pair, not a held-out endpoint protocol or a matched baseline comparison.

3. Roadmap status after this work

Priority	Status	Evidence / remaining work
P0-1 corrected route comparison	DONE - route-level	Paired environment seeds, shared feasibility rules, and action/fallback accounting. The separate demand-driven QKP comparison remains out of scope.
P0-2 physics checks	DONE - internal only	Fiber formula consistency, configured FSO anchor sampling, and temporal correlation. No matched device/channel calibration or measured traces.
P0-3 link sensitivity	DONE - diagnostic	Rate-ratio, QBER-margin, and hard outage-mask tests vary simulator policy inputs only.
P0-4 QBER stress	PARTIAL	Controlled action-probability sweep covers QBER 0.01, 0.03, 0.05, 0.07, 0.09, 0.10, 0.109 and 0.111. The out-of-threshold candidate is masked. Route success, entropy, and preference curves across endpoint/weather conditions remain.
P0-5 FSO stress	PARTIAL	A 0.01x-100x Cn2 input sweep, season/hour quality samples, configured marginal availability, and an outage-mask counterfactual exist. Action-level response only: no route outcomes; pointing/background parameters are fixed and outage-persistence bursts are unsupported.
P0-6 five training seeds	IN PROGRESS - 1/5 complete	Five independent full-budget training runs are not present. Evaluation episodes for one checkpoint do not substitute for training-seed replication.
P1 endpoint / weather holdouts	PARTIAL	The current checkpoint was run on 42 ordered pairs (31/42 success), plus earlier adverse-condition episodes. These are coverage diagnostics, not held-out endpoint/weather tests.
P1 architecture / feature ablations	TODO	No matched-budget GNN depth, attention, feature, or reward ablation set is archived.
P1 inference cost	PARTIAL	Current checkpoint: 42-pair local inference had 3.83 ms median and 4.34 ms p95 decision latency. It is one host/scenario; report a full hardware profile and scaling before deployment claims.
P1 exact small-graph oracle	TODO	No exact feasible-path regret study on generated small graphs is archived.
P2 network-service study	BLOCKED BY MODEL SCOPE	Requires absolute key generation and QKP capacity, time-indexed demand, consumption/refill, blocking, load balance, controller message semantics, prediction baselines, and executable demand-level ILP.
P3 external validation	BLOCKED BY DATA / HARDWARE	Requires matched device/channel measurements or hardware-in-the-loop and separate model-error analysis.

4. Interpretation and next work

The corrected comparison supports a route-level simulator statement only. Report success, successful-route hops, revisits, action-mask violations, and route outcomes; do not infer network capacity or superiority from reward alone. Preserve baseline failures as observed outcomes and state each baseline definition.

Finish the remaining QBER/FSO stress outcomes and the five-seed campaign, then run held-out endpoint/weather experiments and matched-budget ablations. The Johann-style demand/QKP study is a separate extension and depends on absolute key-resource units and a time-indexed demand model. Do not present Johann's BB84 equations as implemented until that model is adopted and validated.

Regression check: 41 tests passed, 0 failed in 68.8 s; 1 warning(s) recorded. Details: `test_results_20261002.json`.

Provenance rule: Label every numeric result as measured, independently calculated, model-sampled, or derived from a policy run. Internal consistency checks do not upgrade an assumption to calibrated physics.